


Online Cardinality Joint Replenishment Problem with Delay

Presented by
Ryder Chen



THE UNIVERSITY OF
SYDNEY





Problem Definition

- What is the Cardinality Joint Replenishment Problem?
- What is Delay?
- What is an Online Algorithm?




Cardinality Joint Replenishment Problem (JRP)

- “Generalisation” of Classical JRP
- Car factory that produces Vans, Trucks and SUVs



- Each car type requires specialised machinery
 - Production cost dominated by machinery hiring cost
 - The cost of hiring n machines is $f(n)$
- e.g. Making 3 Vans and 1 Truck costs $f(2)$


- Customers will request cars from the factory
- Each request will have a deadline
- Problem: The factory must decide when to serve each request whilst minimising costs



- Service 1 costs $f(3)$, Service 2 costs $f(2)$

What is f ?

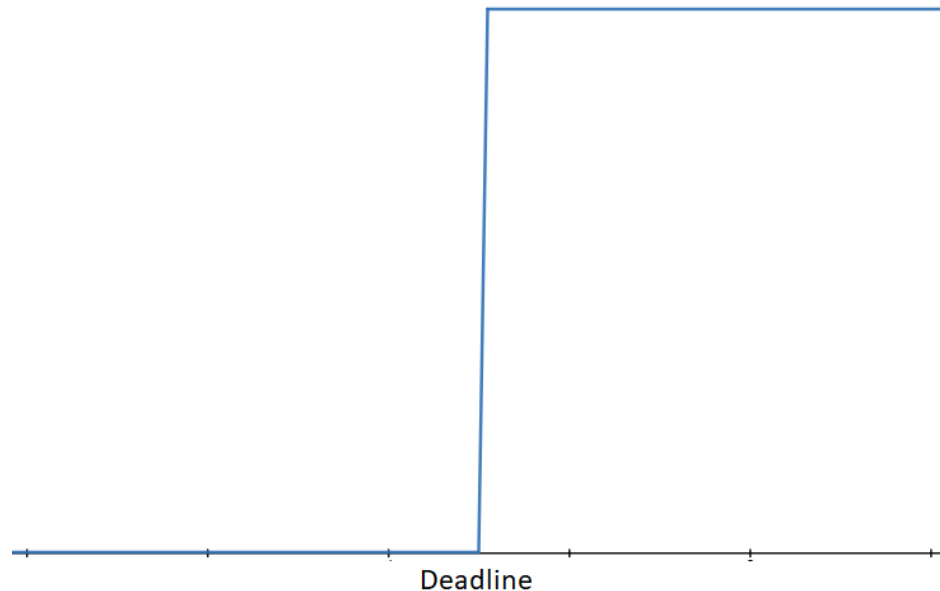
- Non-decreasing concave function



- Every additional piece of machinery hired is successively cheaper

Delay


- Incur a penalty depending on how long a request is left unserved
- Delay generalises deadlines



- Delay is a better representation of reality

Delay


- Each request q has a delay penalty $d_q(t)$
 - arbitrary, non-decreasing, continuous function of time
 - $d_q(t) \rightarrow \infty$ as $t \rightarrow \infty$
- Serving q incurs a penalty of $d_q(t_q)$



Goal: Minimise Ordering + Delay costs

Online Cardinality JRP with Delay

- Offline Algorithm: All information given as input
- Online Algorithm: Runs in real time. Only has information up to the current time.
- Cannot change past decisions.
- No knowledge of the future.




Evaluating Online Algorithms

- Competitive Ratio
- How many times worse is the online solution compared to the optimal, offline solution?

e.g. If $ALG \leq 2 \times OPT$ then the algorithm is 2-competitive OR $O(1)$ -competitive

Background Literature



Offline JRP Literature

Problem	Result	Reference
Classical JRP	NP-Hard	Becchetti et. al. 2009
Classical JRP	1.574-approx. algorithm	Bienkowski et. al. 2013
Cardinality JRP	5-approx. algorithm	Cheung et. al. 2015

Online JRP Literature

	Deadlines	Delay
Classical JRP	3-competitive (Buchbinder et. al. 2008)	3-competitive (Buchbinder et. al. 2008)
Cardinality JRP	$O(1)$ -competitive (Khatkar J. 2020)	

Gap in Literature

- Limited understanding of online JRP
- $O(1)$ -competitive algorithm for Cardinality JRP with Deadlines but what about delay?

	Deadlines	Delay
Classical JRP	3-competitive (Buchbinder et. al. 2008)	3-competitive (Buchbinder et. al. 2008)
Cardinality JRP	$O(1)$ -competitive (Khatkar J. 2020)	?

GOAL: Develop a $O(1)$ -competitive algorithm for Online Cardinality JRP with Delay

Limitations of Existing Work

- Unclear how to generalise Khatkar's algorithm for Cardinality JRP with Deadlines
- Unclear how to generalise other works on JRP and related problems
- Problem is non-trivial and will require innovation and extensions of past work


Results

	Deadlines	Delay
Classical JRP	3-competitive (Buchbinder et. al. 2008)	3-competitive (Buchbinder et. al. 2008)
Cardinality JRP	$O(1)$ -competitive (Khatkar J. 2020)	$O(1)$ -competitive (My Result)

The Algorithm




THE UNIVERSITY OF
SYDNEY



A piecewise reduction

- Approximate $f(x)$ using a piecewise affine function
$$g(x) = \min\{\sigma_l + \delta_l x : l \in [1, n]\}$$
- $f(x) \leq g(x) \leq O(1) \times f(x)$
(Guha et. al. 2001)
- Constant competitive solutions under $g(x)$ are also constant competitive under $f(x)$
- Solution needs to specify which piece of $g(x)$ each service uses as its cost function



Algorithm challenges

1) When do we serve?

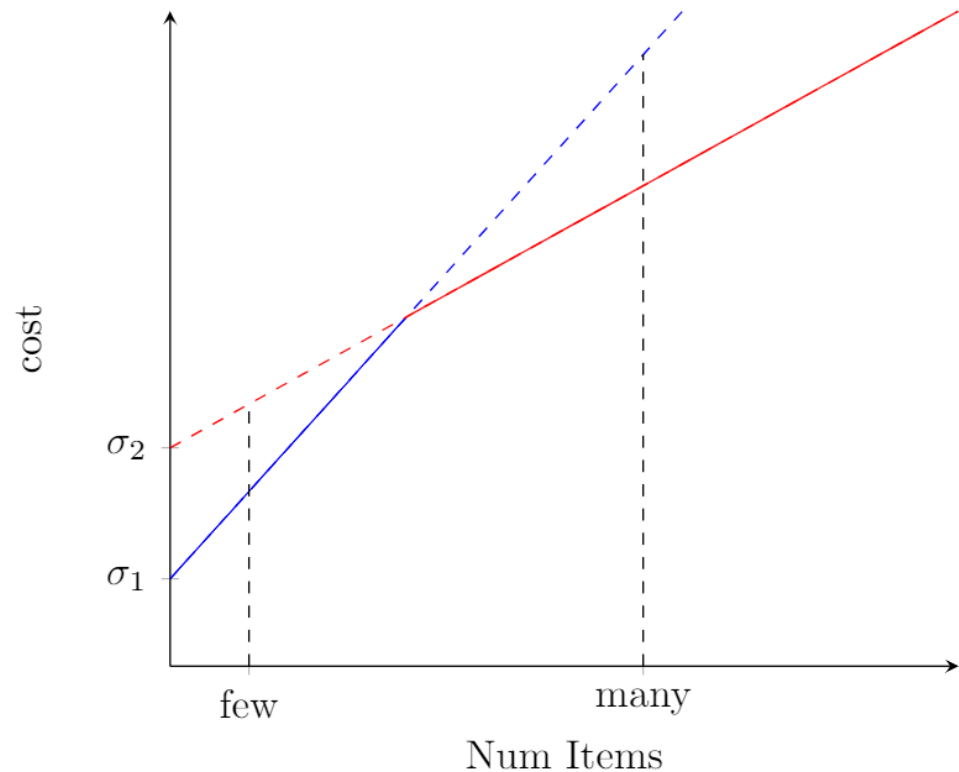
- Want to minimise delay
- No deadlines so how much delay is too much?

2) What do we serve?

- Ordering in bulk is cheaper
- Serving requests on the same item type incurs no extra cost

3) How much to serve?

- Which piece of the piecewise function do we serve with?




Key Contribution 1: Extended Investments

- Traditionally, services pay for the delay on requests it serves
- Investments introduced by Azar and Touitou (2020)
- Instead, pay for delay on requests even on those not served
- Pay for delay accumulated in the future
- Delay costs = Investment costs




Key Contribution 1: Extended Investments

- Azar and Touitou (2020) do not keep track of and use how much has been invested
- **My Innovation:** Keep track of amount invested and serve requests when this reaches a threshold
- This innovation leads to a novel algorithm analysis technique



A specific instance


- Assume the cost function $g(x) = \sigma + \delta x$



- Algorithm does not need to decide how much to serve


When do we serve requests?

- Services will spend at most $O(1) \times \sigma$
- **Idea:** Make a service when delay \approx expected ordering cost
- Wait until requests accumulate total delay of σ




What do we serve?

- Invest a total of σ in the future delay of requests
- If δ is invested in an item type, serve it




What do we serve?

- Invest a total of σ in the future delay of requests
- If δ is invested in an item type, serve it




What do we serve?

- Invest a total of σ in the future delay of requests
- If δ is invested in an item type, serve it




What do we serve?

- Invest a total of σ in the future delay of requests
- If δ is invested in an item type, serve it




What do we serve?

- Invest a total of σ in the future delay of requests
- If δ is invested in an item type, serve it



What do we serve?

- Invest a total of σ in the future delay of requests
- If δ is invested in an item type, serve it



- Balances need to serve as much as possible with leaving requests to wait for other requests on the same item type

Proof sketch

- **Disclaimer:** Very simplified proof sketch
- $ALG = \text{Ordering costs} + \text{Investment costs}$
- $\text{Ordering costs} \leq \text{Investment costs} \leq OPT$
 $\Rightarrow ALG \leq 2 \times OPT$

Ordering costs \leq Investment costs

- Serving a request costs δ
- Serve requests when δ has been invested into it
- An ordering cost of δ is only incurred when an investment cost of δ has already been paid

Investment costs $\leq OPT$


- OPT must serve all requests
- Consider any request q
- There can be at most δ invested in q
- OPT pays δ to serve q
- Investments in $q \leq OPT$'s ordering cost for q
- Apply this argument to all requests to get the desired result

Why is this analysis technique novel?

- Previously: Delay costs \leq Ordering costs $\leq OPT$
- Now: Ordering costs \leq Investment costs $\leq OPT$
- Delay costs = Investment costs
- My approach is the opposite of what is traditionally done


Generalising to arbitrary piecewise affine functions

- Use service levels as done by Khatkar (2020)
- A level l service will serve using the level l piece $\sigma_l + \delta_l x$
- Services will start at level 1 and upgrade over time



Key Contribution 2: Upgrading services

Previously (Khatkar 2020)




Key Contribution 2: Upgrading services

Previously (Khatkar 2020)




Key Contribution 2: Upgrading services

Previously (Khatkar 2020)



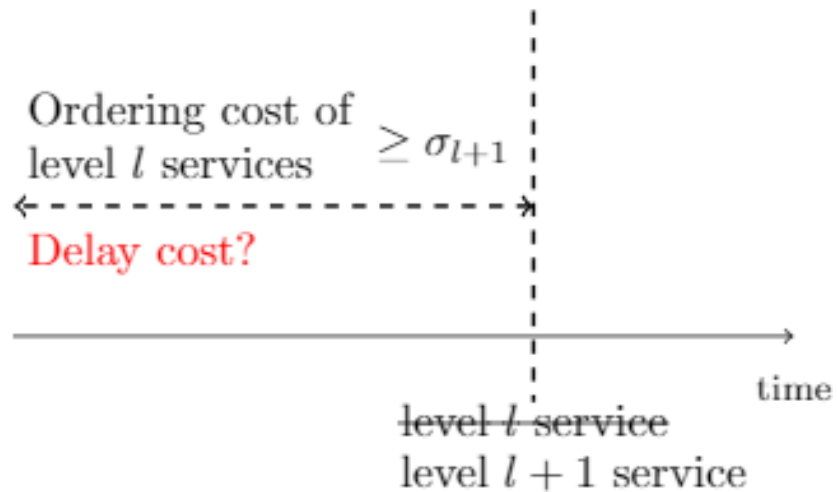
Key Contribution 2: Upgrading services

Previously (Khatkar 2020)



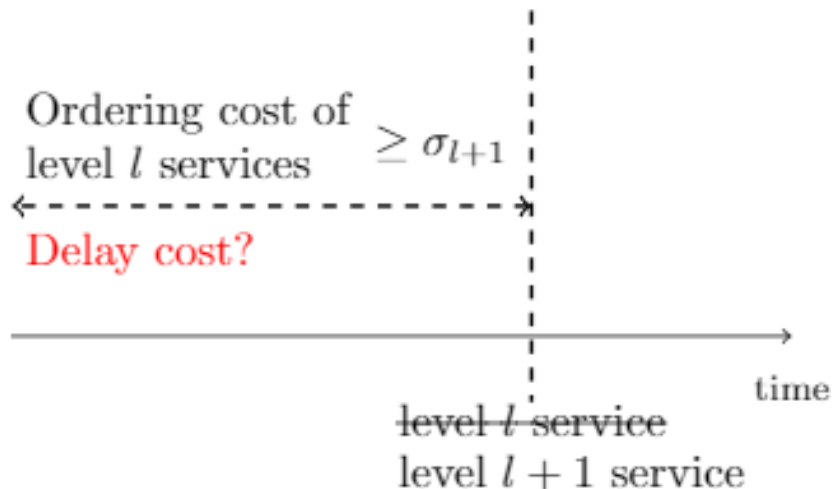
Key Contribution 2: Upgrading services

Previously (Khatkar 2020)

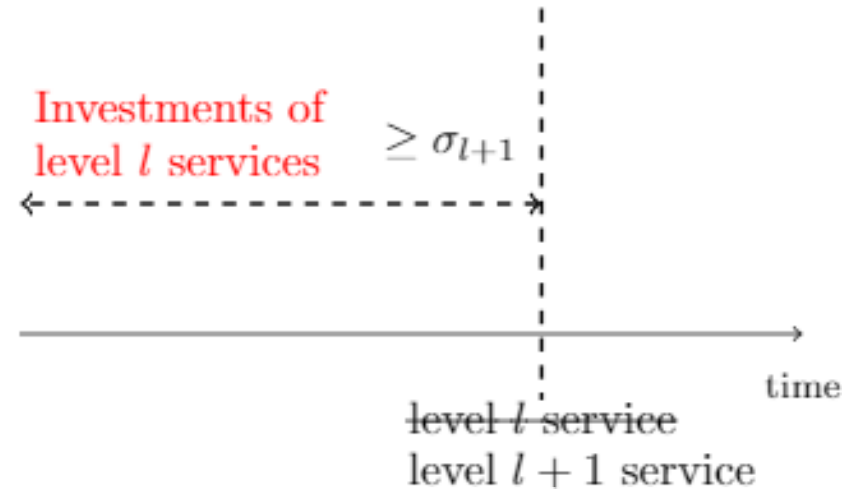


Key Contribution 2: Upgrading services

Previously (Khatkar 2020)



Now (My contribution)




- Ordering costs \leq Investment costs, Investments = Delays
 \Rightarrow My contribution bounds both the ordering and delay costs
- Services made using algorithm from earlier

Summary: Algorithm


1) When do we serve?

Delay \approx Expected Ordering Cost

2) What do we serve?



3) How much to serve?




Summary: Analysis

- $O(1)$ -competitive algorithm for Cardinality JRP with Delay
- Ordering costs \leq Investment costs $\leq OPT$
- $ALG = \text{Ordering costs} + \text{Investment costs}$
 $\Rightarrow ALG \leq 2 \times OPT$



Future Work

- Create online algorithms for other JRP variants
- Applying the idea of extended investments to other problems with delay



Questions?



THE UNIVERSITY OF
SYDNEY



References

- Becchetti, L., Marchetti-Spaccamela, A., Vitaletti, A., Korteweg, P., Skutella, M., Stougie, L.: Latency-constrained aggregation in sensor networks. *ACM Transactions on Algorithms* 6(1), 13:1–13:20 (2009)
- Bienkowski M., Byrka J., Chrobak M., Dobbs N., Nowicki T., Sviridenko M., Swirszcz G., Young N.E. Approximation Algorithms for the Joint Replenishment Problem with Deadlines. In: Fomin F.V., Freivalds R., Kwiatkowska M., Peleg D. (eds) *Automata, Languages, and Programming. ICALP 2013. Lecture Notes in Computer Science*, vol 7965. Springer, Berlin, Heidelberg. (2013)
- Cheung M., Elmachetoub A., Levi R., Shmoys D. The submodular joint replenishment problem. *Mathematical Programming.* 158. 1-27. (2015)
- Buchbinder N., Kimbrel T., Levi R., Makarychev K., Sviridenko M. Online Make-to-Order Joint Replenishment Model: Primal-Dual Competitive Algorithms. *Operations Research.* 61. 952-961. 10.1145/1347082.1347186. (2008)
- Khatkar J., Online Joint Replenishment with Deadline on Subadditive Costs, Honours Thesis, The University of Sydney (2020)
- Azar Y., Touitou N. Beyond Tree Embeddings - a Deterministic Framework for Network Design with Deadlines or Delay. 61st IEEE Annual Symposium on Foundations of Computer Science, FOCS, Durham, NC, USA (2020)
- Guha S., Meyerson A., Munagala K. A constant factor approximation for the single sink edge installation problems. In *Proceedings of the thirty-third annual ACM symposium on Theory of computing (STOC '01)*. Association for Computing Machinery, New York, NY, USA, 383–388. (2001)

Further Proof Details

- $ALG = \text{Shared ordering costs } \sigma + \text{Individual ordering costs } \delta x + \text{Triggering delay costs } \sigma + \text{Investment costs}$
- 3 service types: primary, normal, tail
- Services cost $\leq 3 \sigma$ (excluding individual costs)
- Tail service cost \leq Primary service cost
- Primary services are disjoint so define intervals $[a, t]$ within which we are guaranteed OPT cost \approx primary service cost
- Normal services have an investment cost of σ which we can charge all other costs to. Since Investment costs $\leq OPT$ then normal service costs $\leq OPT$