

Table of Contents

<i>Student Plagiarism: Compliance Statement</i>	<i>i</i>
<i>Acknowledgement</i>	<i>ii</i>
<i>Abstract</i>	<i>iii</i>
<i>Table of Contents</i>	<i>iv</i>
<i>Chapter 1. Introduction</i>	<i>1</i>
1.1 Contribution	2
1.1.1 Publications	3
1.2 Document Outline	3
<i>Chapter 2. Background</i>	<i>4</i>
2.1 Sequencing	4
2.2 After Sequencing – The Computational Pipeline	6
2.3 The need for testing in genomic medicine software	8
2.4 Software testing concepts and approaches	9
2.5 Current testing approaches used in the genomics field	10
2.6 State-of-the-art testing approaches	11
2.6.1 Adaptive Random Testing	11
2.6.2 Metamorphic Testing	12
2.7 Cloud Computing	13
<i>Chapter 3. Aim</i>	<i>17</i>
<i>Chapter 4. Methodology</i>	<i>19</i>
4.1 Testing Framework	19
4.2 Tests	20
4.2.1 Metamorphic Testing	20
4.2.2 Traditional / Simulated Testing	21
<i>Chapter 5. Implementation</i>	<i>23</i>
5.1 Overview	24
5.2 Amazon Web Services – Relevant Background	24
5.2.1 EC2	24
5.2.2 S3	25
5.2.3 Management Console	26
5.2.4 Client and API	26
5.3 Platform and Programming Environment	26
5.4 Tests	27

5.4.1	Metamorphic Tests.....	27
5.4.2	Tests Using Simulated Data	27
5.5	Specific Details of Framework.....	28
5.5.1	AWS Credentials and Security	28
5.5.2	Instance Payment Model and Type	28
5.5.3	AMI Image	28
5.5.4	SSD Option.....	29
5.5.5	Framework Configuration.....	29
5.5.6	GPT Workflow.....	31
5.5.7	GPT Load Balancing and Scalability	31
Chapter 6.	Results	34
6.1	Test Configuration Details	34
6.1.1	AWS	34
6.1.2	Data	35
6.1.3	Details of vGSP under Test.....	36
6.2	Preliminary Testing.....	36
6.2.1	Instance Start-up Times	36
6.2.2	Accessing “local” SSD resources on AWS EC2.....	37
6.3	Main GPT Test	38
6.3.1	Timing Information	39
6.3.2	Cost	42
6.3.3	Test Results.....	44
Chapter 7.	Discussion	49
Chapter 8.	Conclusion.....	53
Appendix 1.	File Type Descriptions	55
8.1	Input Data – FASTQ files	55
8.2	Input Data – Reference Genome – FASTA files	56
8.3	Intermediate output – SAM/BAM files	56
8.4	Output file – VCF	57
Appendix 2.	GPT Configuration File	58
Appendix 3.	Example vGSP Code	59
Appendix 4.	Report File from a GPT Run.....	61
Bibliography	66

CONTENTS

Student Plagiarism: Compliance Statement	ii
Abstract	iii
Acknowledgements	iv
List of Figures	viii
List of Tables	xiii
Chapter 1 Introduction	1
1.1 Background and Motivation	1
1.2 Scientific Contributions and Discoveries	6
1.3 Thesis Structure	7
Part 1. Fungal Classification	8
Chapter 2 Introduction to Fungal Classification	9
2.1 Overview	9
2.2 Fungal Taxonomy	9
2.3 Ribosomal RNA Genes for Classification	10
2.4 Classification Tools for Species Identification	13
2.4.1 Similarity-based Classifiers	13
2.4.2 Phylogeny-based Classifiers	17
2.4.3 Composition-based Classifiers	18
2.4.4 Summary	20
2.5 RDP Naïve Bayes Classifier	22
2.5.1 k -mer Feature Space	22

2.5.2	Statistical-based Classification	23
2.5.3	Bootstrapping	24
Chapter 3	Methods for Fungal Classification	25
3.1	Overview	25
3.2	Fungal ITS Sequence Pre-processing and Curation	25
3.3	Training Set Creation and Evaluation	31
3.4	Validation Set Creation and Classifier Evaluation	32
Chapter 4	Results and Discussion	34
4.1	Overview	34
4.2	Characteristics of the Training Set	34
4.2.1	ITS Sequence Lengths	35
4.2.2	Taxonomic Composition	36
4.3	Training Set Performance	39
4.4	Characteristics of the Validation Set	42
4.4.1	ITS Sequence Lengths	43
4.4.2	Taxonomic Composition	43
4.5	Validation Set Performance	44
4.6	Amplicon Sequencing Simulation Performance	46
4.7	Effects of Sequence Composition on Confidence Values	49
4.8	Future Work	51
4.9	Conclusions	53
Part 2.	Evolutionary Dynamics of Aflatoxin Genes	54
Chapter 5	Introduction to the Evolutionary Dynamics of Aflatoxin Genes	55
5.1	Overview	55
5.2	Aflatoxin Biosynthesis	55
5.2.1	Aflatoxin Types B and G	59
5.3	Evolutionary Dynamics among Fungal Organisms	60
5.3.1	Classification of <i>Aspergillus</i> Species	61

5.3.2	Evolution of <i>Aspergillus</i> Species	63
5.4	Summary	66
Chapter 6	Methods for Analysing Evolutionary Dynamics	67
6.1	Overview	67
6.2	Fungal Genome Assembly	67
6.2.1	Next-Generation Sequencing	68
6.2.2	Velvet Assembler	71
6.2.3	Sequence Error Correction and Assembly of Sequence Reads	76
6.3	Evolutionary Dynamics of Aflatoxin Genes	78
6.3.1	Characterisation of Aflatoxin Gene Pathway	78
6.3.2	Phylogenetic Analysis	80
Chapter 7	Results and Discussion	86
7.1	Overview	86
7.2	Fungal Genome Assembly	86
7.2.1	Assembly Statistics	87
7.2.2	Effects of Coverage Cutoff Parameter	94
7.3	Evolutionary Dynamics of Aflatoxin Genes	95
7.3.1	Characterisation of Aflatoxin Biosynthesis Pathway	95
7.3.2	Phylogenetic Analysis	105
7.4	Future Work	129
7.5	Conclusions	131
Chapter 8	Conclusions	133
	Bibliography	136

CONTENTS

Student Plagiarism: Compliance Statement	ii
Abstract	iii
Acknowledgements	iv
List of Figures	viii
List of Tables	x
Chapter 1 Introduction	1
1.1 Embodied AI	1
1.2 Audio-Visual Navigation	2
1.2.1 Importance of Audio-Visual Navigation	2
1.2.2 Challenges of Audio-Visual Navigation	3
1.3 Contributions	4
1.4 Thesis Outline	5
Chapter 2 Literature Review	6
2.1 Embodied AI in 3D Environments	6
2.1.1 Embodied AI Simulators	6
2.1.2 Embodied AI Tasks	7
2.1.3 Reinforcement Learning	9
2.2 Goal-Oriented Navigation	10
2.2.1 Various Navigation Tasks	11
2.2.2 Navigation Map Construction	14
2.2.3 Memory Mechanism	17
2.3 Audio-Visual Multi-Modality Learning	18
2.3.1 Audio-Visual Representation Learning	19
2.3.2 Embodied Audio-Visual Learning	21

2.3.3	Audio-Visual Navigation	23
Chapter 3	Methods	27
3.1	Input Processing	29
3.2	Visual Mapping	30
3.3	Acoustic Mapping	31
3.4	AFSO-based Audio Encoding	31
3.4.1	Intuition	32
3.4.2	Define Similar Audio Pairs	33
3.4.3	Optimisation Method	38
3.4.4	Batch Sampling	38
3.5	Source Sound Augmentation	41
3.5.1	Audio Mix-up	41
3.5.2	Audio Reverse	42
3.5.3	Sound Augmentation Module	42
3.6	Waypoint Prediction	43
3.7	Path Planning	44
Chapter 4	Experiments and Results	45
4.1	Task Definition	45
4.2	Dataset	46
4.3	Metrics	48
4.4	Implementation	48
4.5	Quantitative Comparisons of Navigation Results	49
4.6	Navigation Results Visualisations	50
4.7	Ablation Studies and Results	54
4.8	Extensive Comparisons on Replica	55
Chapter 5	Discussion	57
5.1	Designs of the AV-GeN Framework	57
5.2	AFSO Method	58
5.2.1	Implications Behind the Contrastive Optimisation	58
5.2.2	Designs of Batch Sampling	59
5.2.3	Designs of Projection Head	60

5.3 Source Sound Augmentation	60
5.4 Limitations in the SoundSpaces AVN Simulator	61
Chapter 6 Conclusion	63
Bibliography	64

CONTENTS

Student Plagiarism: Compliance Statement	ii
Abstract	iii
Acknowledgements	iv
List of Publications	v
List of Figures	ix
List of Tables	xi
Chapter 1 Introduction	1
1.1 Background and Problem Definition	1
1.2 Current Progress and Challenges	1
1.3 Main Contributions	3
1.4 Thesis Outline	4
Chapter 2 Literature Review	5
2.1 Deep Learning Networks and Frameworks	5
2.2 Anomaly Detection in Natural Imaging	7
2.3 Anomaly Detection in Medical Imaging	9
2.4 Memory Networks	11
Chapter 3 Methods	13
3.1 SQUID Overview	13
3.2 Inventing Memory Queue as Dictionary	14
3.2.1 Motivation	14
3.2.2 Space-aware Memory	15
3.2.3 Memory Queue	15
3.3 Gumbel Shrinkage	18

3.3.1	Motivation	18
3.3.2	Combining Hard Shrinkage and Gumbel Softmax	18
3.4	Formulating Anomaly Detection as In-painting	18
3.4.1	Motivation	18
3.4.2	In-painting Block	19
3.4.3	Masked Shortcut	20
3.5	Anomaly Discrimination	20
3.5.1	Motivation	20
3.6	Loss Functions	21
3.7	Network Architectures	22
3.8	Creation of DigitAnatomy	23
Chapter 4	Experiments and Results	26
4.1	Experimental Designs	26
4.1.1	Public Benchmarks	26
4.1.2	DigitAnatomy	27
4.1.3	Comparing Methods	27
4.1.4	Metrics	28
4.1.5	Implementation Details	29
4.2	Image Reconstruction Results on DigitAnatomy	30
4.3	Results on Public Datasets	30
4.4	Ablation Studies	37
4.5	Extensive Studies	38
4.6	Robustness to Disease Samples in the Training Set	42
Chapter 5	Discussion	43
5.1	Limitations	43
5.2	Future Work	44
Chapter 6	Conclusion	47
	Bibliography	49

Contents

Student Plagiarism: Compliance Statement	ii
Abstract	iii
Acknowledgements	iv
List of Figures	viii
Chapter 1 Introduction	1
1.1 Our Contribution	2
1.2 Outline	3
Chapter 2 Preliminary Notions	5
2.1 Spanners	5
2.1.1 The Θ -Graph	6
2.1.2 Bounded Degree and Light Planar Spanners Exist	7
2.2 Online Routing	7
2.2.1 Routing in the Θ -Graph	8
2.2.2 Greedy Routing	8
2.2.3 Compass Routing	9
2.2.4 Right-Hand Routing	9
2.3 Delaunay Triangulations	9
Chapter 3 Everything But Lightness	12
3.1 Description of $\mathcal{BDG}(V)$	12
3.2 Degree Bound of $\mathcal{BDG}(V)$	14
3.3 Spanning Properties of $\mathcal{BDG}(V)$	14
3.3.1 A Separation of Neighbours of Neighbours	14
3.3.2 Notion of Protection	15
3.3.3 Lesser Protection Buys Neighbours in the Same Cone	15

3.3.4	Penultimate and Middle Edges are Bought Full Protection	16
3.3.5	Hull Paths Within Cones are Almost Complete	16
3.3.6	Proof of the Spanning Ratio	17
3.4	Algorithmic Construction of $\mathcal{BDG}(V)$	20
3.5	$\mathcal{BDG}(V)$ Defeats Both Greedy and Compass Routing	21
Chapter 4	Marking the Terrain	22
4.1	Everything is Protected	22
4.2	Signposts	23
4.3	Algorithmic Construction of $\mathcal{MBDG}(V)$	26
Chapter 5	Routing	28
5.1	Chew's Routing Algorithm	28
5.2	Compromised Chew's Routing Algorithm	29
5.2.1	Termination	31
5.2.2	Routing Ratio	31
5.3	Simulating Compromised Chew's Routing Algorithm on $\mathcal{MBDG}(V)$	35
5.3.1	Preliminaries	35
5.3.2	The Simulation	38
Chapter 6	Lightness, and Routing Redux	45
6.1	The Levkopoulos and Lingas Protocol	45
6.1.1	Preliminaries	46
6.1.2	How the Polygon P Grows	46
6.1.3	Condition For Including an Edge	48
6.2	<i>weight</i> Refers to Face Paths	55
6.3	Signposts II	56
6.4	Routing on $\mathcal{LMBDG}(V)$	56
Chapter 7	Conclusion	59
7.1	Further Work	59
7.1.1	Tighter Bounds on the Routing Ratio of $\mathcal{LMBDG}(V)$	60
7.1.2	Navigating Unmarked Terrain	60
7.1.3	Pushing the Degree Down	60
7.1.4	Memoryless Navigators	60

7.1.5	Tighter Weight Bound.....	61
7.1.6	Simulation of Alternative Algorithms.....	61
References		62

CONTENTS

Student Plagiarism: Compliance Statement	ii
Abstract	iii
Acknowledgements	iv
List of Figures	vii
List of Tables	ix
Chapter 1 Introduction	1
Chapter 2 Background	5
2.1 Markov Decision Process (MDP)	5
2.2 Dynamic Programming and Reinforcement Learning	9
2.3 Linear Temporal Logic	13
2.4 Automata	16
Chapter 3 Problem Statement	19
Chapter 4 Related Work	20
4.1 Satisfaction Goals	20
4.2 Challenges of Objective Satisfaction in RL	21
4.3 Objective-based RL for LTL	22
4.3.1 Satisfaction Guarantees without Automata	23
4.3.2 Automata-Based Approaches	23
4.3.3 Miscellaneous RL Settings for LTL	25
4.4 Objective-based RL for Finite LTL	26
4.4.1 Direct Translation to Rewards	26
4.5 Reward Machines	27
4.6 Planning with Finite LTL	29

4.7	RL with LTL Landscape	29
4.8	Finite Horizon Problem in Reward-Based RL	31
4.8.1	Time-Awareness in Deep RL	31
Chapter 5	Methods	32
5.1	Objective-Based RL Theorems	32
5.2	Objective-Based RL Algorithms	36
5.3	Speeding Up the Agent with Heuristics	39
Chapter 6	Experimental Evaluation	41
6.1	Frozen Lake Environment	41
6.2	Reward Machines for Tasks	44
6.3	Optimal Policies	44
6.3.1	Infinite Horizon	45
6.3.2	Finite Horizon	46
6.4	Heuristic Methods	49
6.5	Model Efficiency Experiments	51
6.6	Code	55
Chapter 7	Discussion	56
7.1	Terminal States	56
7.2	Other RL Paradigms	57
7.3	Improving Update Efficiency	57
7.4	Reward Shaping	58
7.5	Time-Dependent and Reward-Machine-Dependent Environment Transitions	59
7.6	Minimising Time	59
Chapter 8	Conclusion	61
	Bibliography	62
	Appendix	66
1	Additional Background	66
2	Additional Figures	66

CONTENTS

Student Plagiarism: Compliance Statement	ii
Abstract	iii
Acknowledgements	iv
List of Figures	vii
Chapter 1 Introduction	1
1.1 Problems	2
1.2 Motivation	2
1.3 Contributions	3
1.4 Related work	6
1.5 Outline	6
Chapter 2 Background	8
2.1 Problem definition	8
2.1.1 Deadline variant	9
2.1.2 Delay variant	9
2.2 Online setting	9
2.2.1 Measuring performance	9
2.2.2 Running time	10
2.2.3 Deterministic vs Randomized	10
2.2.4 Clairvoyance	11
2.3 TCP acknowledgement	12
2.4 Greedy approach	13
2.4.1 Joint replenishment problem	13
Chapter 3 Multi-level aggregation problem	16
3.1 Joint replenishment problem	16

3.1.1	Problem definition	16
3.1.2	Lower bound	16
3.1.3	Deadline variant	18
3.1.4	Delay variant	25
3.2	Multi-level aggregation problem	33
3.2.1	Problem definition	33
3.2.2	Algorithm description	33
3.2.3	Analysis	36
Chapter 4	Set cover with delay problem	46
4.1	Problem definition	46
4.2	Non-clairvoyant setting	46
4.2.1	Intuition	47
4.2.2	Preliminary notations	48
4.2.3	Algorithm description	48
4.2.4	Analysis	51
4.3	Clairvoyant setting	56
4.3.1	Intuition	56
4.3.2	Prize-collecting set cover	57
4.3.3	$O(\log m)$ -competitive algorithm	58
4.3.4	$O(\log n)$ -competitive algorithm	58
Chapter 5	Conclusion	62
	Bibliography	63

CONTENTS

Student Plagiarism: Compliance Statement	ii
Abstract	iii
Acknowledgements	iv
List of Figures	vii
Chapter 1 Introduction	1
1.1 Results and Techniques	4
1.2 Further Related Work	6
1.3 Thesis Outline	10
Chapter 2 Overview of Cardinality JRP Algorithms	11
2.1 Problem Definition: Piecewise Cardinality JRP	11
2.2 A reduction to Piecewise Cardinality JRP	12
2.3 Algorithm Challenges	14
2.4 Algorithm Intuition Overview	15
Chapter 3 Piecewise Cardinality JRP with Deadlines Algorithms	18
3.1 Intuition: Rent-or-Buy JRP with Deadlines	18
3.2 Rent-Or-Buy Cardinality JRP with Deadlines Algorithm	22
3.2.1 Algorithm Description	22
3.2.2 Analysis Sketch	24
3.3 Intuition: Piecewise Cardinality JRP with Deadlines	25
3.4 Piecewise Cardinality JRP with Deadlines Algorithm	28
3.4.1 Algorithm Description	28
3.4.2 Analysis Overview	32
3.4.3 Analysis: Solution Structure	33
3.4.4 Analysis: Charging service costs	33

3.4.5 Analysis: Charging to OPT	37
Chapter 4 Piecewise Cardinality JRP with Delay Algorithm	42
4.1 Intuition	42
4.2 Algorithm Description	47
4.3 Analysis Overview	52
4.4 Analysis: Solution Structure	53
4.5 Analysis: Charging service costs	57
4.6 Analysis: Charging to OPT	60
4.7 Time Complexity Considerations	70
Chapter 5 Weighted Cardinality JRP	74
5.1 Algorithm Description	74
5.2 Analysis	75
Chapter 6 Conclusions and Open Problems	83
Bibliography	84
Appendix A Rent-Or-Buy Cardinality JRP with Deadlines Pseudocode	86
Appendix B Rent-Or-Buy Cardinality JRP with Deadlines Analysis	88
B.1 Analysis: Solution Structure	88
B.2 Analysis: Charging tail service costs	89
B.3 Analysis: Charging to OPT	91

CONTENTS

Abstract	ii
Contents	iii
List of Figures	v
List of Programs	vi
1 Introduction	1
1.1 Contributions	2
2 Background	3
2.1 Data placement	3
2.2 Data migration	5
2.3 Random data layout	6
2.4 Expander graphs	6
2.5 Randomised load balancing	7
3 Analysis	8
3.1 Problem statement	8
3.2 Evaluating a layout	10
3.3 Random data layouts	12
3.4 Prelude to the analysis	13
3.5 Expander graph analysis	14
3.6 Alternate analysis	16
3.7 Baseline model	19
4 Evaluation	21
4.1 The Zipf distribution	21
4.2 Threshold phenomenon	22
4.3 Results	23
4.4 Other sizes	25
5 Partitioning the demand	29
5.1 Analysis: two parts	29
5.2 Partitioning heuristics: two parts	30
5.3 Evaluation: two parts	32

5.4	Analysis: t parts	33
5.5	Partitioning heuristics: t parts	34
5.6	Evaluation: t parts	38
6	Bound on performance	40
6.1	Approximate lower bound	40
6.2	Mixed integer program formulation	42
6.3	Evaluation	44
7	Maximum coverage	46
7.1	Coverage derived from the maximum load	46
7.2	Upper bound	52
7.3	Simulations	55
8	Discussion	58
8.1	Results	58
8.2	Open problems	60
9	Conclusion	62
	Bibliography	64

CONTENTS

Student Plagiarism: Compliance Statement	ii
Abstract	iii
Acknowledgements	iv
List of Figures	ix
List of Tables	xi
Chapter 1 Introduction	1
1.1 Contributions	3
1.1.1 Study of Temporal Compatibility	3
1.1.2 Practical ILP Formalisation	4
1.2 Outline	5
Chapter 2 The Cophylogeny Reconstruction Problem	6
2.1 General Definitions	6
2.1.1 Phylogeny Trees	6
2.1.2 Host-Parasite Associations and Dependence	8
2.1.3 Tanglegrams in Cophylogeny Reconstruction	8
2.2 Problem Definition	9
2.2.1 Events	11
2.2.2 Constraints	18
2.3 Problem Variations	22
2.3.1 Widespread Parasites	22
2.3.2 Non-Independent Parasites	23
2.3.3 Reticulate Networks	23
Chapter 3 Background Work	24
3.1 Early Developments	24

3.1.1	Brooks Parsimony Analysis	24
3.1.2	Reconciled Tree Methods	25
3.2	Toward Contemporary Cophylogeny – Tree Mapping	26
3.2.1	Event Costs	27
3.2.2	Pareto Sets	28
3.3	Current Heuristics and Approaches	28
3.3.1	Cost Matrix	29
3.3.2	Jungles	29
3.3.3	TARZAN	30
3.3.4	JANE	30
3.3.5	CORE-PA	31
3.3.6	TREEMAP 3	32
3.4	Intractibility and Approximations Results	32
3.4.1	Moving Back Landing Site Problem	32
3.4.2	The Cophylogeny Reconstruction Problem	33
3.5	Summary	33
Chapter 4	Temporal Incompatibilities and Host Switch Resolution	35
4.1	Background	36
4.1.1	Anatomy of a Host Switch	36
4.1.2	Incompatibilities	37
4.1.3	Newer Approaches	38
4.2	Feasible Reconstructions and Event Orders	39
4.2.1	Phylogeny Implicit Orders	39
4.2.2	Host Switch Orders	40
4.3	j -vertex – The Jungle Method	44
4.3.1	Sufficiency	45
4.3.2	Complexity	46
4.4	Proposed Method	47
4.4.1	Necessity	48
4.4.2	Sufficiency	50
4.4.3	Complexity	53
4.5	Summary	54

Chapter 5 Integer Linear Program Formulation	55
5.1 Previous ILP Formulation	55
5.1.1 Fixed Event Times	56
5.2 Proposed ILP Formulation	57
5.2.1 Model Assumptions	58
5.2.2 ILP Formulation	59
5.3 Formulation Correctness	66
5.3.1 Necessity	66
5.3.2 Sufficiency	69
5.3.3 Event Counts	71
5.4 Variable and Constraint Counts	73
5.5 Summary	73
Chapter 6 Reference Implementation	74
6.1 Implementations	74
6.1.1 JANE	74
6.1.2 Proposed Integer Linear Program (ILP)	74
6.1.3 Previous ILP	75
6.2 Experiments	75
6.2.1 Platform	75
6.2.2 Datasets	76
6.3 Results	78
6.3.1 Synthetic Datasets	78
6.3.2 Real World Datasets	81
Chapter 7 Conclusion	84
7.1 Future Work	84
7.1.1 Pushing the bounds of intractability further	84
7.1.2 Expanding the model of reconstruction	85
7.2 Conclusions	86
Appendix A Corrections to Previous ILP Formulation	87
A.1 Missing Constraints	87
A.2 Incorrect Constraints	88

Appendix B	Full ILP Listing	91
Appendix C	Full Result Tables	93
C.1	Column Descriptions.....	93
C.2	Synthetic Datasets.....	94
C.3	Real World Datasets.....	103
Bibliography		105

CONTENTS

Student Plagiarism: Compliance Statement	ii
Abstract	iii
Acknowledgements	iv
List of Figures	vii
List of Tables	ix
Notation	x
Chapter 1 Introduction	1
1.1 Contribution	2
1.2 Outline	3
Chapter 2 Background	4
2.1 Approximation & Randomised Algorithms	4
2.2 Set Function Optimisation Problems	5
2.2.1 Matroid Constraints	6
2.2.2 Knapsack Constraints	7
2.2.3 Independent Set in Interval Graphs	7
2.3 Algorithm Techniques	7
2.3.1 Greedy Algorithms	8
2.3.2 Continuous Greedy	9
2.3.3 Contention Resolution Schemes	12
2.3.4 Pipage Rounding	14
Chapter 3 Preliminaries	17
3.1 Maximum Facility Location	17
3.2 Integer Program	18

3.3 Linear Relaxation.....	19
Chapter 4 Algorithms	20
4.1 Independent Rounding	20
4.1.1 Algorithm Engineering.....	24
4.2 Dependent Rounding	25
4.3 Hardness.....	32
4.3.1 Reducing from Maximum Coverage	32
Chapter 5 Conflicting Predictions from Multi-mapping Reads	34
5.1 Modelling the Problem	35
5.2 Implementation	36
5.3 Related Work.....	38
5.4 Experiments.....	39
5.4.1 Simulated Reads: Craig Venter	39
5.4.2 Illumina Reads: NA12878.....	43
Chapter 6 Discussion	47
6.1 Algorithms	47
6.2 Experiments.....	48
6.3 Open Problems and Future Work	48
Chapter 7 Conclusion	50
Bibliography	52

CONTENTS

Student Plagiarism: Compliance Statement	iii
Abstract	iv
Acknowledgements	v
List of Figures	viii
List of Tables	x
Chapter 1 Introduction	1
Chapter 2 Background	5
2.1 The Changing Nature of Database Systems	5
2.1.1 The Issues with Mutual Exclusion as a Synchronisation Tool	6
2.2 Non-blocking Synchronisation Techniques	7
2.2.1 Definition of Non-Blocking	7
2.2.2 Atomic Primitives	7
2.2.3 The ABA Problem	9
2.2.4 Classifying Primitives	9
2.2.5 Multi-word Primitives	10
2.2.6 Non-Blocking Data Structures	11
Chapter 3 Problem Statement and Methodology	14
3.1 Implementing a Non-Blocking Key-Value Store	14
3.1.1 The Key-Value Store Interface	15
3.1.2 Non-Blocking Synchronisation	15
3.1.3 Linearisability: A Correctness Condition	15
3.1.4 Memory Reclamation for Unmanaged Languages	16
3.2 Testing Performance in the Face of Contention	16
3.2.1 Synchrobench	16

3.2.2	The Comparison Skip List: Fraser's Non-Blocking Skip List	17
3.2.3	Definitions of Performance Metrics	17
3.2.4	Experimental Settings	17
3.2.5	The Test Machine	18
Chapter 4	The Rotating Skip List	19
4.1	Overview of the No Hot Spot Non-Blocking Skip List	19
4.2	The Naive Design	20
4.3	Consolidation of Index Nodes into Rotating Arrays	21
4.4	Reducing the Frequency of Background Maintenance	24
4.5	Delegating Physical Deletes to the Background Thread	26
4.6	Memory Reclamation for The Rotating Skip List	29
4.6.1	Reducing the Frequency of Memory Reclamation	31
Chapter 5	Discussion and Suggestions for Future Work	33
5.1	Insights Reached During Development of The Rotating Skip List	33
5.2	Suggestions for Future Work	34
5.2.1	Improving Performance in Other Test Scenarios	34
5.2.2	Expanding the Interface of The Rotating Skip List	36
Chapter 6	Conclusion	38
	Bibliography	39